


CS4833: Embedded Systems

Instruction Set Architecture (ISA) and ARM processor

CS 4833: Embedded Systems



Reviews

- Basic components: hardware
 - Transistors and gates
 - Combinational and sequential logic circuits
 - Timers, counters etc.
- General hardware architectures
 - von Neumann vs. Harvard
- Processing engines in embedded systems
 - Microprocessor, microcomputer, microcontroller
 - Digital signal processors (DSPs)
 - Graphics processing units (GPUs)
- Understand hardware words
 - Data, address or instructions

CS 4833: Embedded Systems

2



Outline

- Overview of ISA: outside view
 - Types: arithmetic/logic, data transfer & flow control
 - Data transfer and address modes
- Register Transfer Language (RTL): inside view
 - Registers & register operations: read & write
 - Concept of data-path
 - Instruction cycle: fetch, decode, execute, next
- Example: a simple ISA & processor design
- Case study: ISA for ARM processors

Instruction Set for A Microprocessor

- Instruction set specifies basic operations supported by one microprocessor
 - **Outside view** of what can be done: ability
 - Operations: transfer/store data, operate data and make decisions
- Different types of instructions
 - Data transfer: move/exchange data with components
 - Arithmetic and logic: computational capabilities/functions
 - Flow of control: order of instructions to be executed

Data Transfer Instructions

■ Typical data transfer instructions

➤ **assembly** vs. machine codes

LD destination, source	Load—source operand transferred to destination operand can be either register or memory location.
ST source, destination	Store—source operand transferred to destination operand source must be a register and the destination must be memory.
MOVE destination, source	Transfer from register to register or memory to memory.
XCH destination, source	Interchange the source and destination operands.
PUSH/POP	Operand pushed onto or popped off of the stack.
IN/OUT destination, source	Transfer data from or to an input/output port.


Data Transfer Instructions (cont.)

■ Pieces information are needed

➤ Where is the **data: source**

➤ Where to go: **destination**

■ Possible locations of data: register, memory, I/O






Addressing Modes: More about Operands

Where is the data?

- How to understand source/destination

Addressing modes

- Immediate
- Direct and Indirect
- Register direct/indirect
- Indexed
- Program counter relative



CS 4833: Embedded Systems




Addressing Mode: Immediate

Typical instructions

STI	- Store immediate
LDI / LOADI	- Load Immediate
MOVI	- Move Immediate

Destination maybe implicit

Limitation on value of immediate data



int x = 0xB;

MOVE OPR1, #BH

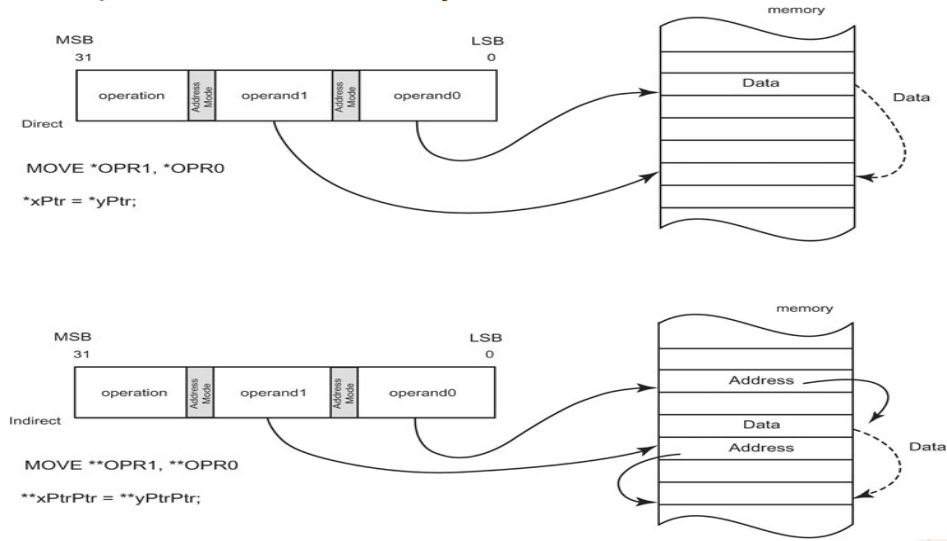
CS 4833: Embedded Systems

8



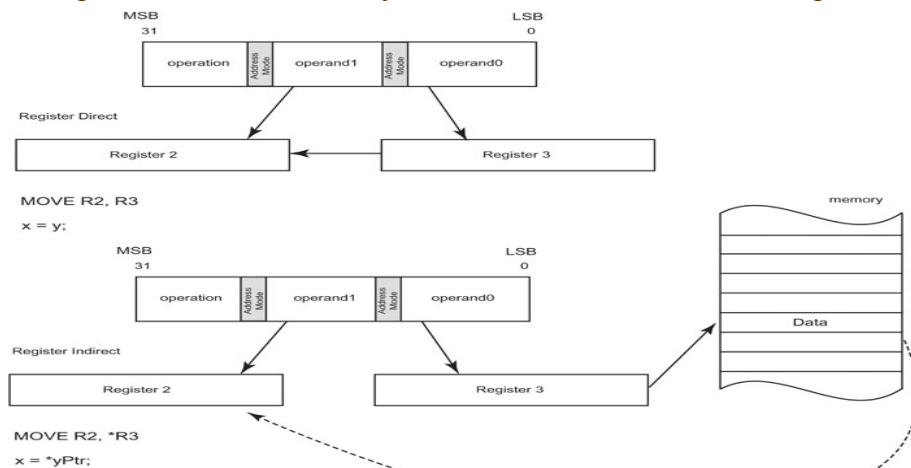
Addressing Mode: Direct/Indirect

- Operand is the memory address of needed data

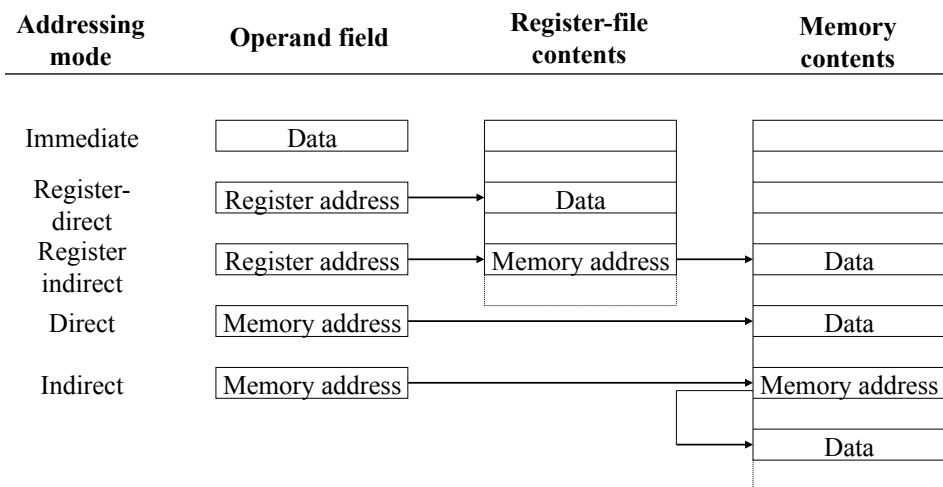


Addressing Mode: Register Direct/Indirect

- Register direct: data in register operands
- Register indirect: memory location of needed data in registers



Addressing Modes



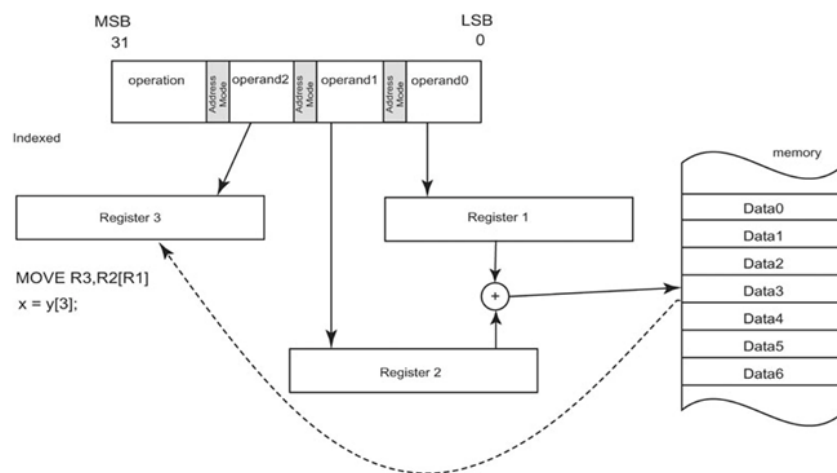
CS 4833: Embedded Systems

11



Addressing Mode: Indexed

- Base and offset of data items → access array



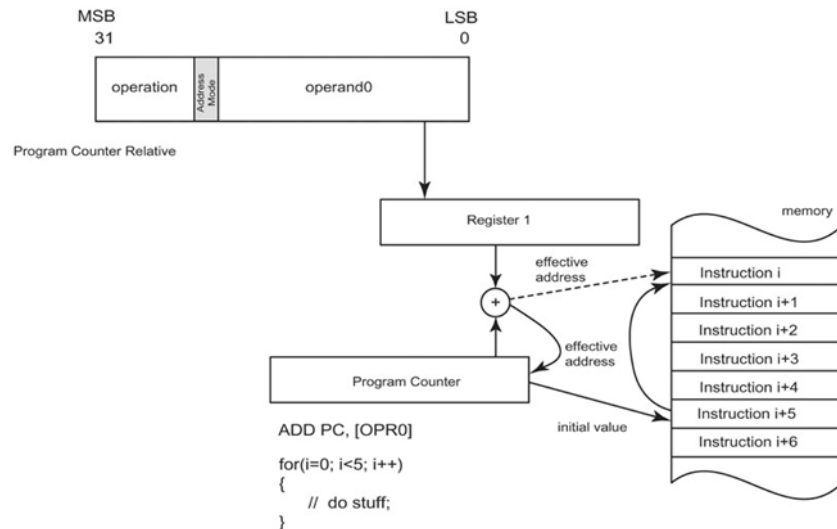
CS 4833: Embedded Systems

12

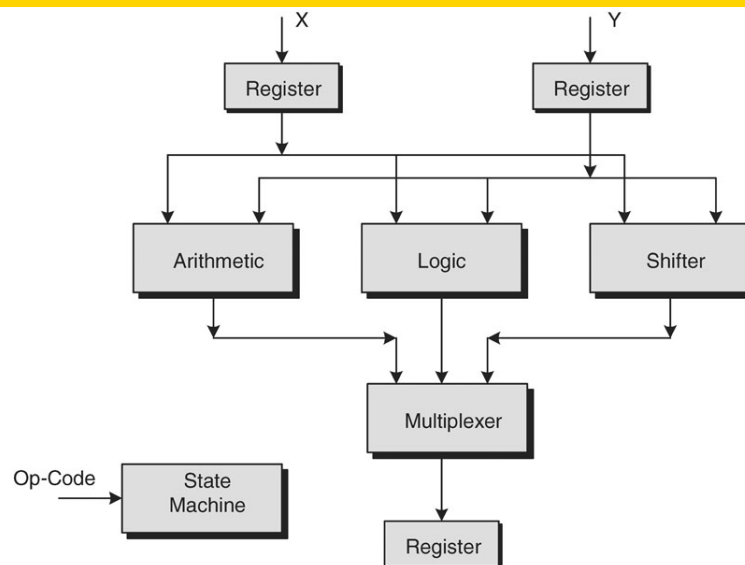


Addressing Mode: PC Relative

- For flow of control → access instructions



Arithmetic and Logic Instructions: ALU



CS 4833: Embedded Systems

14



Arithmetic and Logic Instructions

Typical instructions

ADD2, ADD3	// Two or three operands addition
ADDC	addition with carry
SUB2, SUB3	// Two or three operands subtraction
SUBB	subtraction with borrow
MUL	multiplication
DIV	division
INC	increment
DEC	decrement
TEST	operand tested and specified condition set
TESTSET	atomic test and set

AND	bitwise AND
OR	bitwise OR
XOR	bitwise Exclusive OR
NOT or INV	complement
CLR or SET	clear or set
CLRC, SETC	carry manipulation

SHR operand, count	logical shift right
SHL operand, count	logical shift left
SHRA operand, count	arithmetic shift right
SHLA operand, count	arithmetic shift left
ROR operand, count	rotate right
ROL operand, count	rotate left

CS 4833: Embedded Systems

15



Execution Flows

Sequential

- Instructions are executed sequentially

Branch

- Decision point and instructions from different parts of program continues

Loop

- Repeatedly execute a set of instructions

Procedure/function call

- Need to leave current *running context*

CS 4833: Embedded Systems

16



Branch Instructions

■ Branch conditions and instructions

E, NE	Operand1 is <i>equal/not equal</i> to Operand2.
Z, NZ	The result of the operation is <i>zero/not zero</i> .
GT, GE	Operand1 is <i>greater than/greater than or equal</i> to Operand2.
LT, LE	Operand1 is <i>less than/less than or equal</i> to Operand2.
V	The operation resulted in an <i>overflow</i> —the result is larger than can be held in the destination.
C, NC	The operation produced a <i>carry/no carry</i> .
N	The result of the operation is <i>negative</i> .

BR label	unconditional branch to the specified label
BE label, BNE label	branch to the specified label if the equal flag is set or not set
BZ label, BNZ label	branch to the specified label if the zero flag is set or not set
BGT label	branch to the specified label if the greater than flag is set
BV label	branch to the specified label if the overflow flag is set
BC label, BNC label	branch to the specified label if the carry flag is set or not set
BN label	branch to the specified label if the negative flag is set

CS 4833: Embedded Systems


17



Procedure/Function Call

■ Structure and instructions

■ Running context may change: save running states



CALL operand	PC is unconditionally saved and replaced by specified operand; control is transferred to specified memory location.
RET	Previously saved contents of PC are restored, and control is returned to previous context.


CS 4833: Embedded Systems

18



Microprocessor: Register View

- **Internal view:** how things are got done (how signals are moved/changed within core **hardware**)
- Registers: containing data/address etc.
- **Control and data-path**
- Basic registers
 - Parallel
 - Shifted
- Register operations
 - Read
 - Write

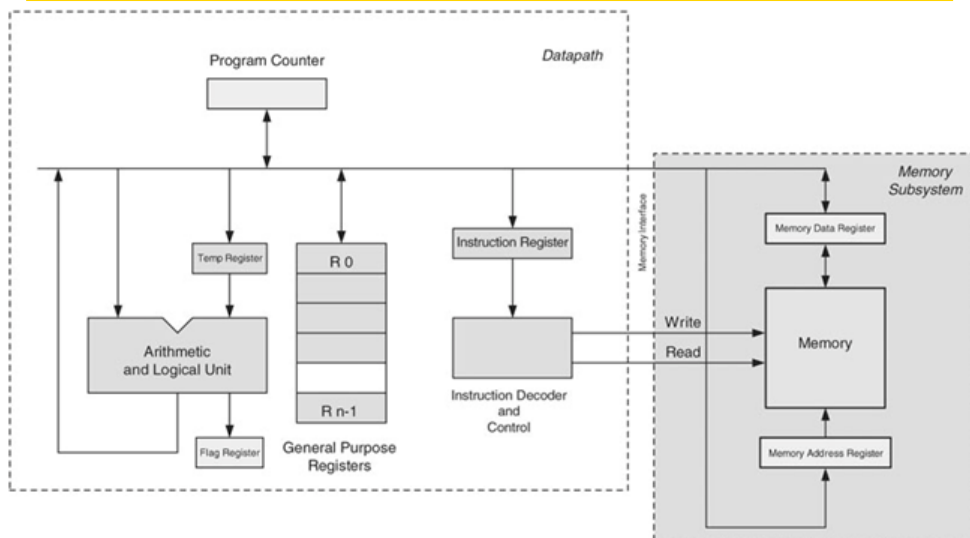


CS 4833: Embedded Systems

19



Closer Look: Register View



CS 4833: Embedded Systems

20



Register Transfer Language (RTL)

- Describe how data is transferred: registers/memory
- Can easily translate to HDL or Verilog VHDL for implementation

Operator	Operations
\leftarrow	Transfer from right hand side to left hand side
\rightarrow	If-then operation
[index]	Select word from memory at index

$+ - \times \div$	Arithmetic operation


Instruction and Register Transfer Notation

Type	Instruction	ISA level	Register Transfer
Data transfer	Move Register	MOVE R1, R2	$R1 \leftarrow R2$
	Move from memory	MOVE R1, memadx	$R1 \leftarrow (\text{memadx})$
	Move to memory	MOVE memadx, R1	$(\text{memadx}) \leftarrow R1$
	Move immediate	MOVE R1, #DEAD	$R1 \leftarrow \#DEAD$
Control Flow	Unconditional branch	BR \$1	$PC \leftarrow \$1$
	Conditional branch	BNE \$1	$\text{Cond}(PC \leftarrow \$1)$ $\text{if}(\text{cond}) PC = \1
Logic	Complement accum	CMA	$A \leftarrow \neg A$
	AND register	AND R1, R2	$R1 \leftarrow R1 \wedge R2$
	Shift register	SHL R1, #3	$R1 \langle 3..0 \rangle \leftarrow R1 \langle 31-n..0 \rangle \#(n@0)$
Arithmetic	ADD register w. carry	ADDC R1, R2	$R1 \leftarrow R1 + R2 + C$
	Clear carry	CLRC	$C \leftarrow 0$
Prog. control	Donot execute inst.	NOP	$PC \leftarrow (PC+1)$
	Stop execution	HALT	$PC \leftarrow PC$

Instruction Cycle

Fundamental operations/steps when executing instructions

- Fetch: retrieve instruction from memory
- Decode: analyze op-code
- **Mem**: get data from memory
- Execute: perform operation
- Next: get next instruction addr. → PC



CS 4833: Embedded Systems

Register View for the Steps

Fetch

```
MAR ← PC;           // PC enabled out to bus, MAR captures value
MDR ← Memory[MAR];  // contents of specified memory location placed into MDR
IR ← MDR;           // MDR enabled out to bus, IR captures value
```

Execute

```
// C Instruction
*xPtr = y;
// ISA Level Instruction
ST *R1, R2;
// RTL Level Instructions
MAR ← R1;           // R1 enabled out to bus, MAR captures value
MDR ← Memory[MAR];  // contents of specified memory location placed into MDR
R2 ← MDR;           // MDR enabled out to bus, R2 captures value
```

CS 4833: Embedded Systems

24



Example: A Simple (Trivial) Instruction Set

Assembly instruct.	First byte		Second byte	Operation
MOV Rn, direct	0000	Rn	direct	$Rn = M(\text{direct})$
MOV direct, Rn	0001	Rn	direct	$M(\text{direct}) = Rn$
MOV @Rn, Rm	0010	Rn	Rm	$M(Rn) = Rm$
MOV Rn, #immed.	0011	Rn	immediate	$Rn = \text{immediate}$
ADD Rn, Rm	0100	Rn	Rm	$Rn = Rn + Rm$
SUB Rn, Rm	0101	Rn	Rm	$Rn = Rn - Rm$
JZ Rn, relative	0110	Rn	relative	$PC = PC + \text{relative}$ (only if Rn is 0)
	opcode		operands	

CS 4833: Embedded Systems

25



An Example of Compilation

C program

```
int total = 0;
for (int i=10; i!=0; i--)
    total += i;
// next instructions...
```

compilation →

Equivalent assembly program

```
0  MOV R0, #0;      // total = 0
1  MOV R1, #10;     // i=10
2  MOV R2, #1;      // constant 1
3  MOV R3, #0;      // constant 0
Loop: JZ R1, Next;   // Done if i=0
5  ADD R0, R1;      // total += i
6  SUB R1, R2;      // i--
7  JZ R3, Loop;     // Jump always
Next: //next instruction ...
```

CS 4833: Embedded Systems

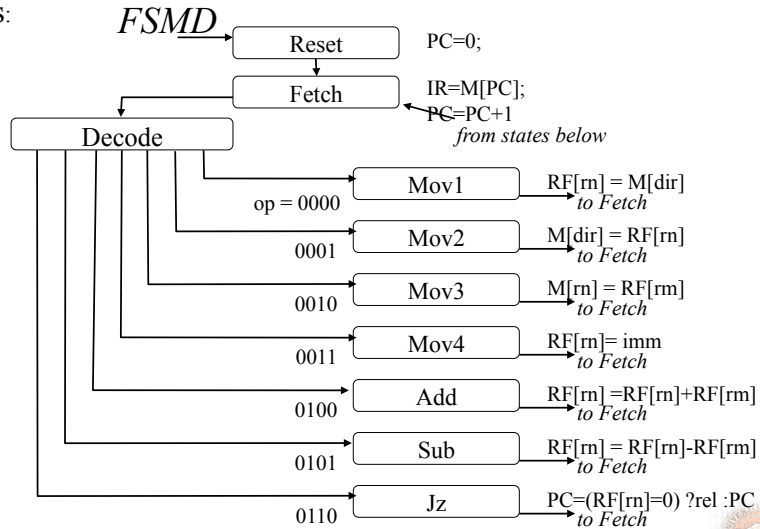
26



Designing a General Purpose Processor

Declarations:

bit PC[16];
bit IR[16];
bit M[64k][16],
bit RF[16][16];

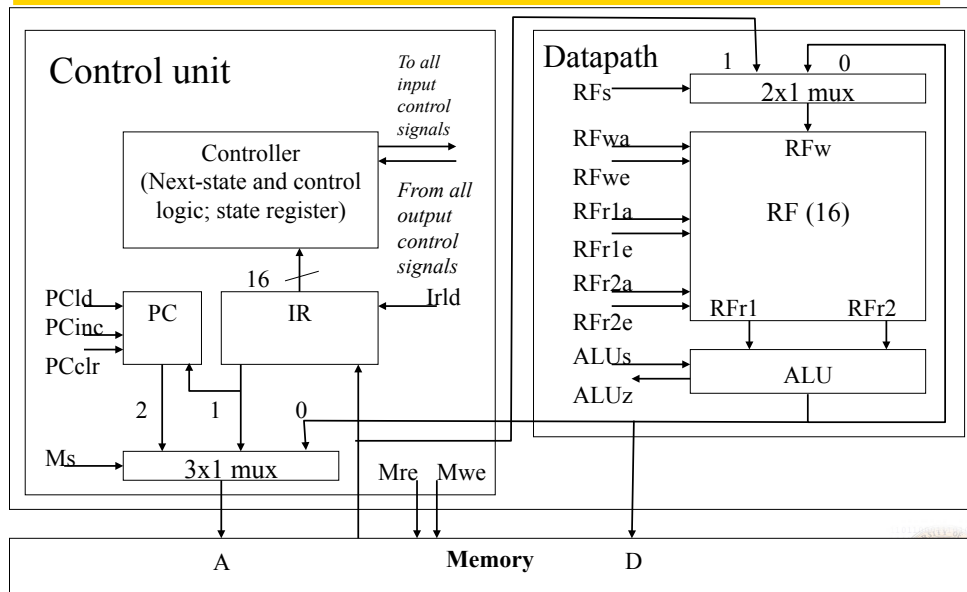


CS 4833: Embedded Systems

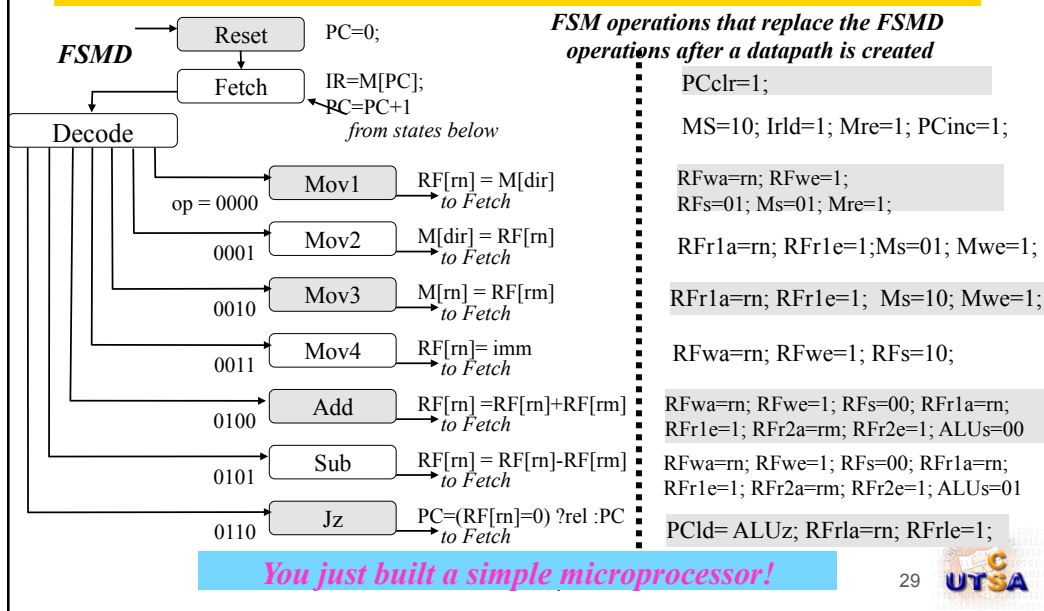
27



Architecture of a Simple Microprocessor



The Controller of the Microprocessor



Case Study: ARM Instruction Set

- ARM assembly language – RISC
- ARM programming model
- ARM memory organization
- ARM data operations (32 bits)
- ARM flow of control
- Hardware-based floating point unit

ARM Overview

■ Versions of ARM

- Several versions have been extended
- Focus on ARMv7 and ARMv8

■ Adopt standard assembly language

- Examples: load (LDR) and ADD

```
                LDR r0,[r8] ; a comment
label          ADD r4,r0,r1
```

ARM Data Types

■ ARM address: 32 bits long


- Address refers to byte

■ Word: 32 bits long

- Can be divided into four 8-bit bytes

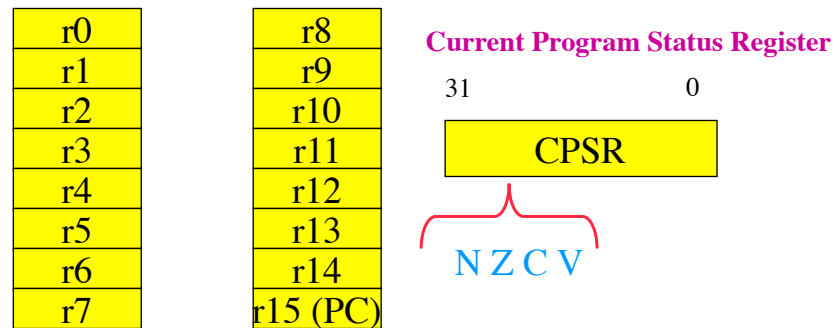
■ Endianness: configured at power-up

- Little vs. big endian mode
- Bytes vs. endianness



ARM Programming Model

■ Registers



CS 4833: Embedded Systems

33



ARM Status Bits

■ Every arithmetic, logical, or shifting operation sets CPSR bits:

- N (negative)
- Z (zero)
- C (carry)
- V (overflow)

■ Examples: note that $-1 = 0xffffffff$ (two's-complement)

- $-1 + 1 = 0xffffffff + 0x1 = 0$: NZCV = 0110.
- $(2^{31}-1)+1 = 0x7fffffff + 0x1 = 0x80000000 = -2^{31} \rightarrow$ NZCV = 1001.

CS 4833: Embedded Systems

34



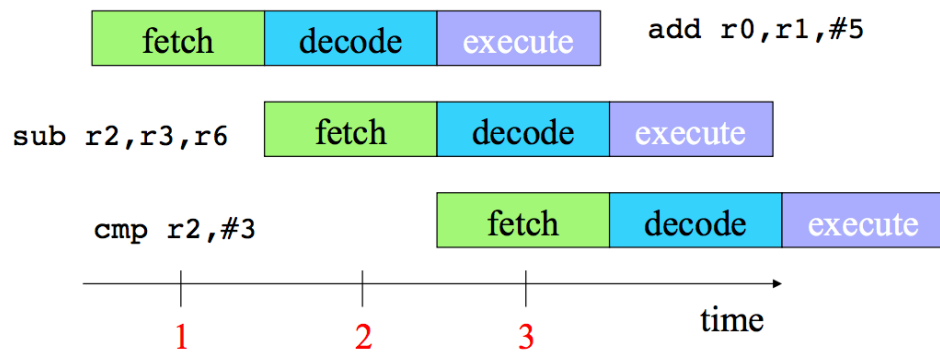
ARM Basic Instructions and Pipeline

- Basic format: `ADD r0,r1,r2`

- Computes $r1+r2$, stores in $r0$.

- Immediate operand: `ADD r0,r1,#5`

- Computes $r1+5$, stores in $r0$.



ARM Data Instructions

- ADD, ADC : add (w. carry)
- SUB, SBC : subtract (w. carry)
- RSB, RSC : reverse subtract (w. carry)
- MUL, MLA : multiply (and accumulate)
- AND, ORR, EOR
- BIC : bit clear
- LSL, LSR : logical shift left/right
- ASL, ASR : arithmetic shift left/right
- ROR : rotate right
- RRX : rotate right extended with C

ARM Comparison Instructions

- CMP : compare
- CMN : negated compare
- TST : bit-wise test
- TEQ : bit-wise negated test
- These instructions set only the NZCV bits of CPSR.

ARM load/store/move Instructions

- LDR, LDRH, LDRB : load (half-word, byte)
- STR, STRH, STRB : store (half-word, byte)
- Addressing modes:
 - register indirect : LDR r0, [r1]
 - with second register : LDR r0, [r1,-r2]
 - with constant : LDR r0, [r1, #4]
- MOV, MVN : move (negated)
MOV r0, r1 ; sets r0 to r1

Other Addressing Modes

- Base-plus-offset addressing:

```
LDR r0,[r1,#16]
```

- Loads from location r1+16

- Auto-indexing increments base register:

```
LDR r0,[r1,#16]!
```

- Post-indexing fetches, then does offset:

```
LDR r0,[r1],#16
```

- Loads r0 from r1, then adds 16 to r1.

Example: Assignment in C (1)

- C code: $x = (a+b)-c$;

- Assembler:

```
ADR r4,a      ; get address for a
LDR r0,[r4]   ; get value of a
ADR r4,b      ; get address for b, reusing r4
LDR r1,[r4]   ; get value of b
ADD r3,r0,r1  ; compute a+b
ADR r4,c      ; get address for c
LDR r2,[r4]   ; get value of c
SUB r3,r3,r2  ; complete computation of x
ADR r4,x      ; get address for x
STR r3,[r4]   ; store value of x
```

Example: Assignment in C (2)

■ C code: $z = (a \ll 2) | (b \& 15)$

■ Assembler:

```
ADR r4,a      ; get address for a
LDR r0,[r4]   ; get value of a
MOV r0,r0,LSL 2 ; perform shift
ADR r4,b      ; get address for b
LDR r1,[r4]   ; get value of b
AND r1,r1,#15 ; perform AND
ORR r1,r0,r1  ; perform OR
ADR r4,z      ; get address for z
STR r1,[r4]   ; store value for z
```

ARM Flow Control

■ All operations can be performed conditionally, by testing CPSR:

➤ EQ, NE, CS, CC, MI, PL, VS, VC, HI, LS, GE, LT, GT, LE

■ Branch operation: B #100

➤ Can be performed conditionally

ARM subroutine linkage

- Branch and link instruction:

BL foo

- Copies current PC to r14.

- To return from subroutine:

MOV r15, r14

ARM ISA Summary

- Load/store architecture

- Most instructions are RISC

- operate in single cycle.
- Some multi-register operations take longer.

- All instructions can be executed conditionally

- ARMv7

- Deployed in Cortex-A15 (Snapdragon Krait, Nvidia Tegra, TI OMAP), Cortex-A5 (AMD Fusion), Atmel microcontrollers

ARMv8: Next Generation

- Addition of 64-bit support
 - Larger virtual address space
- New instruction set (A64)
 - Fewer conditional instructions
 - New instructions to support 64-bit operands
 - No arbitrary length load/store multiple instructions
- Enhanced cryptography (both 32 and 64-bit)
- Mostly backwards compatible with ARMv7
- Expansion into higher performance markets
 - Mobile phones, servers, supercomputers
 - Cortex-A53, Apple Cyclone, Nvidia Denver

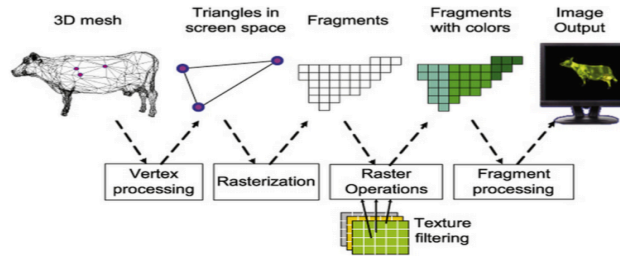
CS 4833: Embedded Systems

45



Graphics Processing Units (GPUs)

- Specialized processors for images to display
 - Primary elements: vertex and pixel shaders
- Compute Unified Device Architecture (CUDA)
 - GPGPUs: used for general purpose processing
 - Normally hundreds of simple cores
- Single program multiple data (SPMD)
 - Threading directives and memory addressing



46



GPU Programming

- OpenCL: general language for GPU, CPU & DSP

- Open standard with emphasis on portability
- Supported by Intel, AMD, Nvidia and ARM

- CUDA (Compute Unified Device Architecture)

- Enable general purpose GPU (GPGPU)
- Specific to NVIDIA with SPMD

```
void saxpy_serial( int n,
                  float a,
                  float *x,
                  float *y)
{
    for (int i=0; i<n; i++)
        y[i] = a*x[i] + y[i];
}

// perform SAXPY on 1M elements
saxpy_serial(4096*256, 2.0, x, y);
```



```
global
void saxpy_parallel(int n,
                   float a,
                   float *x,
                   float *y)
{
    int i = blockIdx.x * blockDim.x +
          threadIdx.x;
    if (i<n) y[i] = a*x[i] + y[i];
}

// perform SAXPY on 1M elements
saxpy_parallel<<<4096,256>>>(n,2.0, x, y);
```

CS 4833: Embedded Systems

47



GPU Programming (cont.)


- Kernel: scales across any number of parallel cores

- Contains grid of thread blocks

- Thread block: 1D, 2D, 3D

- All threads in a block run the same code, over same data in shared memory space
- Threads have ID numbers within block to select work and address shared data
- Threads in different blocks cannot cooperate

- Threads are grouped by warps (scheduling units)



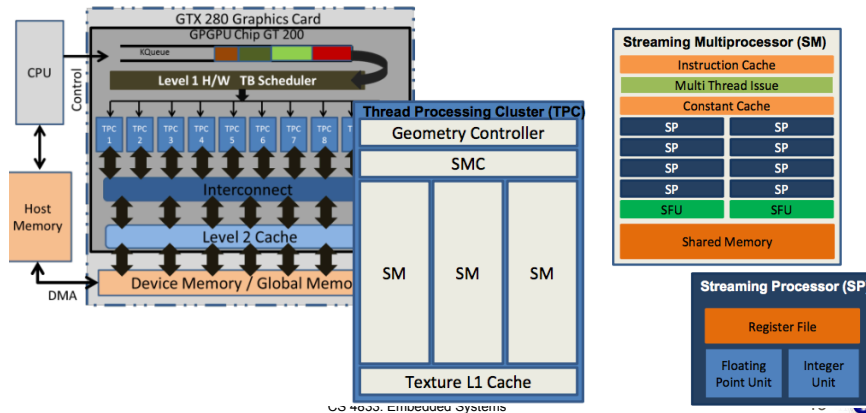
CS 4833: Embedded Systems

48



Nvidia GeForce Block Diagram

- Ultra low power (ULP) version in Tegra 3 & 4
- Desktop GTX 280 (embedded version w. diff. number of TPUs, SMs. etc.)



RP3: VideoCore IV GPU

- GPU: VideoCore IV dual-core @ 400MHz
- Low-power mobile multimedia processor architecture
- Highly integrated w. CPU, memory and display
- High power efficiency in driving fast off-chip buses
- Encoder and decoder support
 - Hardware block assisting JPEG
 - 1080p30 Full HD H.264 Video
- Runs common libraries and drivers
 - OpenGL ES 2.0, OpenVG 1.1, OpenMAX
 - ✓ OpenGL: cross-language, cross-platform application programming interface (API) for rendering 2D and 3D vector graphics

CS 4833: Embedded Systems

50



VideoCore IV GPU Architecture

GPU core: parallel computing of video @ low clock speed

- Co-processor extension to ARM
- Grouped by slices: shared common cache

Quad Processor Unit (QPU): 4-way SIMD machine

- Four-way physical parallelism
- 16-way virtual parallelism (4-way over four successive clock cycles)
- RP3: 12 QPUs

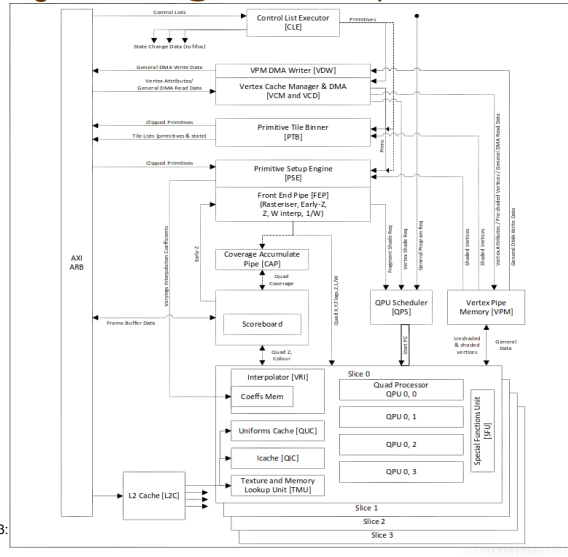
Special function Units (SPU)

- SQRT, LOG, RECIPSQRT, EXP

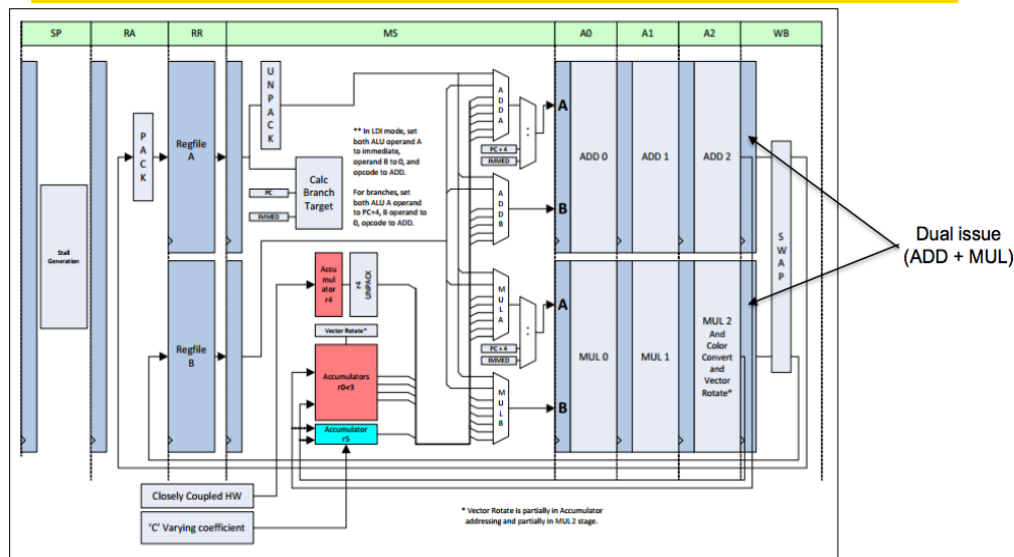
Texture and memory units

- Specialized memory for 3D texture data

CS 4833:



VideoCore IV GPU Pipeline: 8 Stages



CS 4833: Embedded Systems

52 UTSA

Summary

- Overview of ISA: outside view
 - Types: arithmetic/logic, data transfer & flow control
 - Data transfer and address modes
- Register Transfer Language (RTL): inside view
 - Registers & register operations: read & write
 - Concept of data-path
 - Instruction cycle: fetch, decode, execute, next
- Example: a simple ISA & processor design
- Case study: ISA for ARM processors